

Local Party Websites in Germany: A Chapter-Year Panel, 2015–2025

Data report for the public v1.0.0 release

Hanno Hilbig

2026-07-11

Table of contents

Summary	1
Background and Summary	2
Methods	2
Sampling frame and website discovery	2
Archival reconstruction	3
Content extraction and classification	3
From records to panel measures	3
Data Records	3
Technical Validation	4
Human validation of the extraction layers	4
Issue attention tracks the national agenda	4
Campaign content tracks the electoral calendar	6
Composition and platform gradients	6
Migration attention and the local far right	6
Coverage and selection into observability	8
Usage Notes	9
Code Availability	10
Citation	10
References	10

Summary

What do local party organizations do and say? Answering this question at scale has been difficult because the activities of local chapters are rarely digitized in any systematic form. We introduce a chapter-year panel that measures the observable web presence of 2,611 local chapters (predominantly county-level Kreisverbände) of the seven main German parties (CDU, CSU, SPD, Grüne, Linke, AfD, FDP) in municipalities with at least 5,000 residents, observed annually from 2015 to 2025 (28,721 chapter-years). Drawing on roughly one million archived pages from the Internet Archive’s Wayback Machine, we reconstruct each chapter’s website in each year and extract structured measures: leadership boards and their composition, issue emphasis, digital and mobilization infrastructure, and a classified news layer of 269,548 archived articles labeled by policy topic, article function, and geographic focus. All extraction and classification ran on a locally hosted open-weight language model, and every model-generated layer was validated by a human coder, with agreement on 288 of

288 reviewed records. The public release contains chapter-year aggregates and coverage metadata only: no personal data and no reproduced website text. The data should be read as a panel of observable local party web presence, not as a census of local party activity.

Background and Summary

Local party organizations recruit members, select candidates, and connect national politics to local communities. While a rich literature studies national and state parties through manifestos, speeches, and press releases, the local level has remained largely out of empirical reach: local manifestos and personnel lists are rarely archived, and most studies substitute national party positions for local ones. This dataset takes a different route. It measures what local chapters put on their own websites, and it does so longitudinally by reconstructing each site from archived snapshots.

The dataset makes three contributions. First, it provides the first systematic panel of German local party web presence, built against a fixed sampling frame rather than an ad hoc collection of sites, so that coverage and its limits are measurable. Second, it extracts comparable organizational measures across parties and years: board composition, issue emphasis crosswalked to the Comparative Agendas Project topic scheme, and the mobilization functions of the party-website literature (Gibson and Ward 2000). Third, it adds a classified news layer that captures what chapters communicate month to month, which supports analyses of local issue attention and campaign activity over time.

Two design decisions shape everything downstream. The unit of observation is the chapter-year: for each chapter and each year from 2015 to 2025, we select archived snapshots of the chapter's known URL and extract measures from the pages captured in that year. And because whether a chapter-year is observable depends on website existence and archival coverage, every row carries explicit coverage metadata. We return to the resulting selection issues in the validation and usage sections; they are the single most important thing to understand before using the data.

Methods

Sampling frame and website discovery

The frame of possible chapters is constructed before any website discovery. We restrict the Destatis Gemeindeverzeichnis (the official municipality register) to municipalities with at least 5,000 residents, which yields 2,943 municipalities aggregating to 399 Kreise. Crossing these Kreise with the parties' organizational scope rules (AfD, FDP, Grüne, Linke, and SPD organize everywhere; the CDU outside Bavaria; the CSU only in Bavaria) produces 2,394 party-by-Kreis cells. Each cell is a possible chapter, and all coverage statistics in the release are computed against this frame.

We then locate chapter websites per cell through four channels, in decreasing order of authority: (i) the parties' own chapter directories (999 chapters); (ii) the SPD's published organizational roster (8); (iii) systematic probing of predictable domain patterns such as `cdu-kreisname.de` (869); and (iv) search-engine queries with model-classified results (735), where queries contained only party and place names. The channel is recorded per chapter in the `url_source` variable, so users can condition on discovery quality. Discovery found a site for 97.8 percent of cells, and the unfound cells concentrate in the parties with the weakest local digital presence. One caveat deserves emphasis: discovery ran against the live web in early 2026, so chapters whose sites went offline earlier without replacement are missing entirely. Early panel years therefore carry survivorship selection on top of archival selection, a limitation we disclose but have not yet quantified.

Archival reconstruction

Drawing on the Wayback Machine’s CDX index, we query each chapter’s known URL for every year, select snapshots under a structural-page-priority rule, and fetch the archived HTML, roughly 1.03 million pages in total. Pages are classified by function (homepage, officials, news, issues, contact, and so on), and the per-chapter-year snapshot selection is recorded in `chapter_year_provenance.csv` so any aggregate can be audited. Whether a chapter-year is observable at all thus depends on two processes we do not control: the chapter maintaining a site, and the Internet Archive having captured it that year. Rather than treating this as a nuisance, we measure it: a deterministic observability ladder (`observability_status`) classifies every chapter-year, and a release coverage taxonomy assigns each row a basis, bucket, and tier.

Content extraction and classification

All content extraction and classification ran on a locally hosted open-weight language model (Qwen3-30B-A3B-Instruct) with schema-constrained decoding, which forces outputs into the target record structure. No scraped page content was sent to commercial cloud APIs. From the archived pages we extract officials with roles, issue content blocks, and organizational contact information. News pages are deduplicated to one record per unique article (269,548 articles), and each article is classified along three dimensions: a policy topic family, an article function (press release, event announcement, position statement, election campaigning, or council work), and a geographic focus (local, regional, or national).

Three derived classifications complete the record layer. Gender is inferred from first names, a standard method benchmarked at roughly 94 percent accuracy (Santamaría and Mihaljević 2018); the inference is binary, unknown names are retained and flagged, and `gender_coverage` reports the classified share per chapter-year. Issue topics are assigned by a dictionary with a model fallback into about 25 families, crosswalked to Comparative Agendas Project major topics. Role strings are normalized into board-role categories. These classifiers also ran locally, so no model-generated value in the release originates from a commercial API.

From records to panel measures

A research layer aggregates the extracted records into deterministic chapter-year measures; no free text survives into the release. Every substantive feature carries a companion status column that distinguishes an observed positive value, an observed zero, coverage that exists but was not extracted, and missing coverage. This four-way distinction is the core of the release design: an apparent zero is never conflated with a missing site or a missing archive.

Data Records

The release contains eight data files at two grains, one row per chapter or one row per chapter-year, with identifiers stable across files.

Table 1: Release data files.

File	Grain	Rows	Cols	Size
<code>chapter_roster.csv</code>	chapter	2,611	15	0.4 MB
<code>chapter_year_panel.csv</code>	chapter-year	28,721	42	13.3 MB
<code>chapter_year_panel_enriched.csv</code>	chapter-year	28,721	71	7.3 MB

File	Grain	Rows	Cols	Size
<code>chapter_year_topic.csv</code>	chapter-year-family	25,832	5	1.0 MB
<code>chapter_year_news.csv</code>	chapter-year	13,883	13	0.8 MB
<code>chapter_year_news_topic.csv</code>	chapter-year-family	85,238	5	3.6 MB
<code>chapter_year_provenance.csv</code>	chapter-year	28,721	19	2.0 MB
<code>topic_family_cap_crosswalk.csv</code>	family	26	4	small

The base panel carries identity and geography (party, organizational level, municipality, Kreis, the `ags` municipality key, population), the coverage block (`observability_status`, the coverage taxonomy columns, and scope flags), and eleven substantive features with their status columns: board listing, official counts, issue block counts, an office address flag, six chapter-specific social-platform flags with national party accounts filtered out, and the deduplicated news article count. The enriched panel adds derived measures: board gender composition, chair presence and gender, turnover, professionalization proxies, an `officials_source` flag identifying whether the board was drawn from party-office or council pages, web-activity counts, seven mobilization link flags following the Gibson-Ward functions, and fifteen wide issue-topic presence flags. The news file reports per chapter-year article volume together with label shares by geographic focus and article function, and the two long-format topic files carry per-family counts for issue blocks and news articles respectively, each joined to its CAP major topic. Full variable definitions are in `CODEBOOK.md`.

Technical Validation

Human validation of the extraction layers

Every model-generated layer was validated by a human coder reading extracted values next to source-page excerpts. We drew four stratified samples: (i) a structured-extraction spot check of 48 records across the officials, issues, and about routes; (ii) 63 news articles checked field by field; (iii) 100 issue topic-family assignments; and (iv) 77 news articles with all three content labels. Across all 288 reviewed records the coder agreed with the extraction. The public release reports aggregate agreement counts in `coverage_and_quality/validation_human_agreement.md`; the underlying sheets contain source excerpts and remain private.

Issue attention tracks the national agenda

Beyond record-level accuracy, we ask whether the classified news layer recovers known movements of the German political agenda. It does. Figure 1 plots the share of news topics devoted to four issue families. Attention to integration and migration peaks at 9.5 percent of topics in 2015, the year of the refugee crisis, and collapses thereafter. Environment and climate attention builds through the climate-protest wave and peaks in 2019, the year of the “climate election” to the European Parliament. Health jumps from around 2 percent to over 8 percent in the pandemic years 2020 and 2021, and energy spikes in 2022, the year of the gas crisis. None of these patterns is available to the classifier, which sees article text without dates. Local party communication, it turns out, follows the national agenda closely.




Figure 1: Issue attention in local party news tracks national agenda shocks. Each panel shows one topic family's share of all classified news topics by year; shaded bands mark the relevant national event.

Campaign content tracks the electoral calendar

The article-function labels pass an analogous check. Figure 2 shows the share of articles classified as election campaigning: it peaks in federal and European election years and bottoms out in off-years, from 4.8 percent in 2015 to 13.4 percent in 2024. Again, the classifier does not know the year.




Figure 2: Campaign-content share of archived local party news articles by year. BTW marks federal elections, EP European Parliament elections.

Composition and platform gradients

The board-composition measures reproduce known party differences. Figure 3 orders parties by their mean board female share: the Grünen, with a parity quota, sit at 47 percent; Linke and SPD, with formal quotas, follow; the CDU and CSU, with informal quorums, sit near 28 percent; and the FDP and AfD, without quotas, trail at 23 and 16 percent. The ordering matches the parties' candidate-quota rules exactly, although we stress that the shares are name-inferred and website-listed rather than administrative counts.

Platform adoption shows an equally sharp gradient. Telegram presence is essentially an AfD phenomenon: it appears after 2018, accelerates in 2021 when mainstream platforms tightened moderation, and reaches 15 percent of AfD chapters by 2025, while remaining near zero for all other parties (Figure 4). Both patterns emerge from the extraction pipeline without any party-specific rules.

Migration attention and the local far right

Two further patterns connect the data to active debates on how mainstream parties respond to the radical right. The accommodation literature asks whether mainstream parties win voters back by adopting radical-right issue positions, and most evidence suggests they do not (Krause, Cohen and Abou-Chadi 2023). Our data show what the underlying issue competition looks like at the local level. Figure 5 plots the migration share of each party's news topics: the AfD leads every mainstream party in every single year, mainstream attention collapses after 2016 rather than rising




Figure 3: Mean board female share by party, chapter-years with at least 70 percent of board names classified. Party labels note the applicable gender quota rule.




Figure 4: Share of chapters linking a Telegram channel, AfD versus all other parties.

to meet the challenger, and only the AfD re-spikes during the 2023 migration debate. At least on their own websites, local mainstream chapters disengaged from the migration issue; we find little evidence of accommodation from below.




Figure 5: Attention to integration and migration by party, as a share of each party’s classified news topics.

The second pattern speaks to the debate on far-right normalization and local rootedness (Schroeder, Ziblatt and Bochart 2025). Figure 6 compares the geographic focus of AfD news with that of all other parties. The AfD is by a wide margin the most nationally oriented party on the ground: 39 percent of its article content was nationally focused in 2015, against 6 to 18 percent for the others. While that gap persists, the AfD’s local-focus share climbs from 48 percent in 2015 to 73 percent in 2025, with visible interruptions in the federal election years 2021 and 2024. Taken together, the two panels show a party that entered local politics as a national protest platform and progressively localized its communication: organizational rooting, measured directly from what chapters publish.

Coverage and selection into observability

Of 28,721 chapter-years, 25,245 are in analysis scope; the remainder is out-of-universe or quarantined after roster-integrity checks. The release denominator is researcher-usable coverage: 16,866 chapter-years (66.8 percent of scope), of which 14,199 meet the full structured-coverage threshold (`coverage_basis = full`) and 2,667 carry observed feature values from below-threshold or fallback-routed pages (`coverage_basis = feature_level`). Most of the remainder is irreducible: 7,083 chapter-years (28.1 percent) have no same-year archived snapshot.

Coverage is strongly non-random, and Figure 7 shows where. Full-threshold coverage rises from about 35 percent in 2015 to 73 percent in 2025, and party levels range from 74.7 percent (SPD) to 37.2 percent (AfD). The CSU panel is the most volatile: many CSU chapters live on shared




Figure 6: Geographic focus of news articles, AfD versus all other parties.

state-party infrastructure whose archival capture varies year to year. A descriptive linear-probability model in `missingness_diagnostics.csv` shows usable coverage rising with municipality population and over time. The implication for users is direct: missing cells must not be treated as zeros, and analyses should condition on `observability_status` or `coverage_basis`.

Usage Notes

Three usage rules prevent the most likely mistakes. First, read `ags` and `kreis_id` as strings: both are zero-padded and silently lose their leading zero under default integer inference, which affects roughly 80 percent of chapters. Second, filter to `analysis_scope_included = 1`, and then either to `coverage_basis = 'full'` for a conservative sample or to any non-blank `coverage_basis` for the full researcher-usable denominator, with the basis as a robustness split. Third, treat a feature value as informative only when its status column reports an observed positive or an observed zero.

Several measurement caveats follow from the construction. The extracted boards pool internal party officers, electoral candidates, and elected council representatives; `officials_source` flags the dominant source, and board size, gender, and completeness measures should be conditioned on it. Issue fields measure salience rather than position. Link and social-media flags reflect a site's template, and presence means a link appeared on an archived snapshot, not that the account was active. The news count reflects Wayback archival density as well as posting volume, and the news label shares are composition measures conditional on posting. A single annual snapshot may fall before or after a local election or a site redesign, so year-to-year change is noisy; the data are more reliable for broad cross-sectional comparisons than for fine annual dynamics. Finally, chapters are matched to one seat municipality, which simplifies a multi-municipality organizational geography.

The public tier ships aggregates and coverage metadata only. Person-level records and reproduced page text are withheld for data-protection reasons and are available to researchers on request under a use agreement. Public data and documentation are licensed CC-BY-4.0; code is MIT.




Figure 7: Usable structured coverage by party and year. Each panel shows the share of analysis-scope chapter-years meeting the full coverage threshold.

Code Availability

The full pipeline is available in the code repository (github.com/hhilbig/local-party-websites-germany). Extraction prompts and schemas are in `src/extraction/`, applied verbatim by the local-inference extractors in `scripts/hive/`; the classifier label sets are in `src/enrichment/classifiers.py`; the issue dictionary and CAP crosswalk are in `config/`. The news layer is built by scripts 62 through 65 and labeled by scripts 72 and 73. Scripts 20, 21, 55, and 56 build the chapter-year layers, script 52 the coverage taxonomy, and script 57 regenerates the release bundle deterministically. Human validation was coded in the review application at `tools/validate/`, and the release figures are reproduced by `release/figures_report.py` from the public CSVs alone.

Citation

Hilbig, Hanno (2026). *Local Party Websites in Germany: A Chapter-Year Panel, 2015–2025*, Version 1.0.0. DOI pending Zenodo deposit.

References

- Gibson, R. and Ward, S. (2000). A proposed methodology for studying the function and effectiveness of party and candidate web sites. *Social Science Computer Review* 18(3), 301–319.
- Krause, W., Cohen, D. and Abou-Chadi, T. (2023). Does accommodation work? Mainstream party strategies and the success of radical right parties. *Political Science Research and Methods* 11(1), 172–179.
- Schroeder, V., Ziblatt, D. and Bochert, F. (2025). Is the Brandmauer breaking from below?

- Party cooperation with the AfD in German local councils. WZB Discussion Paper.
- Comparative Agendas Project. Master codebook (21 major topics).
 - Santamaría, L. and Mihaljević, H. (2018). Comparison and benchmark of name-to-gender inference services. *PeerJ Computer Science* 4:e156.